

CS201 – Lecture 9

IA32 Procedures

RAOUL RIVAS

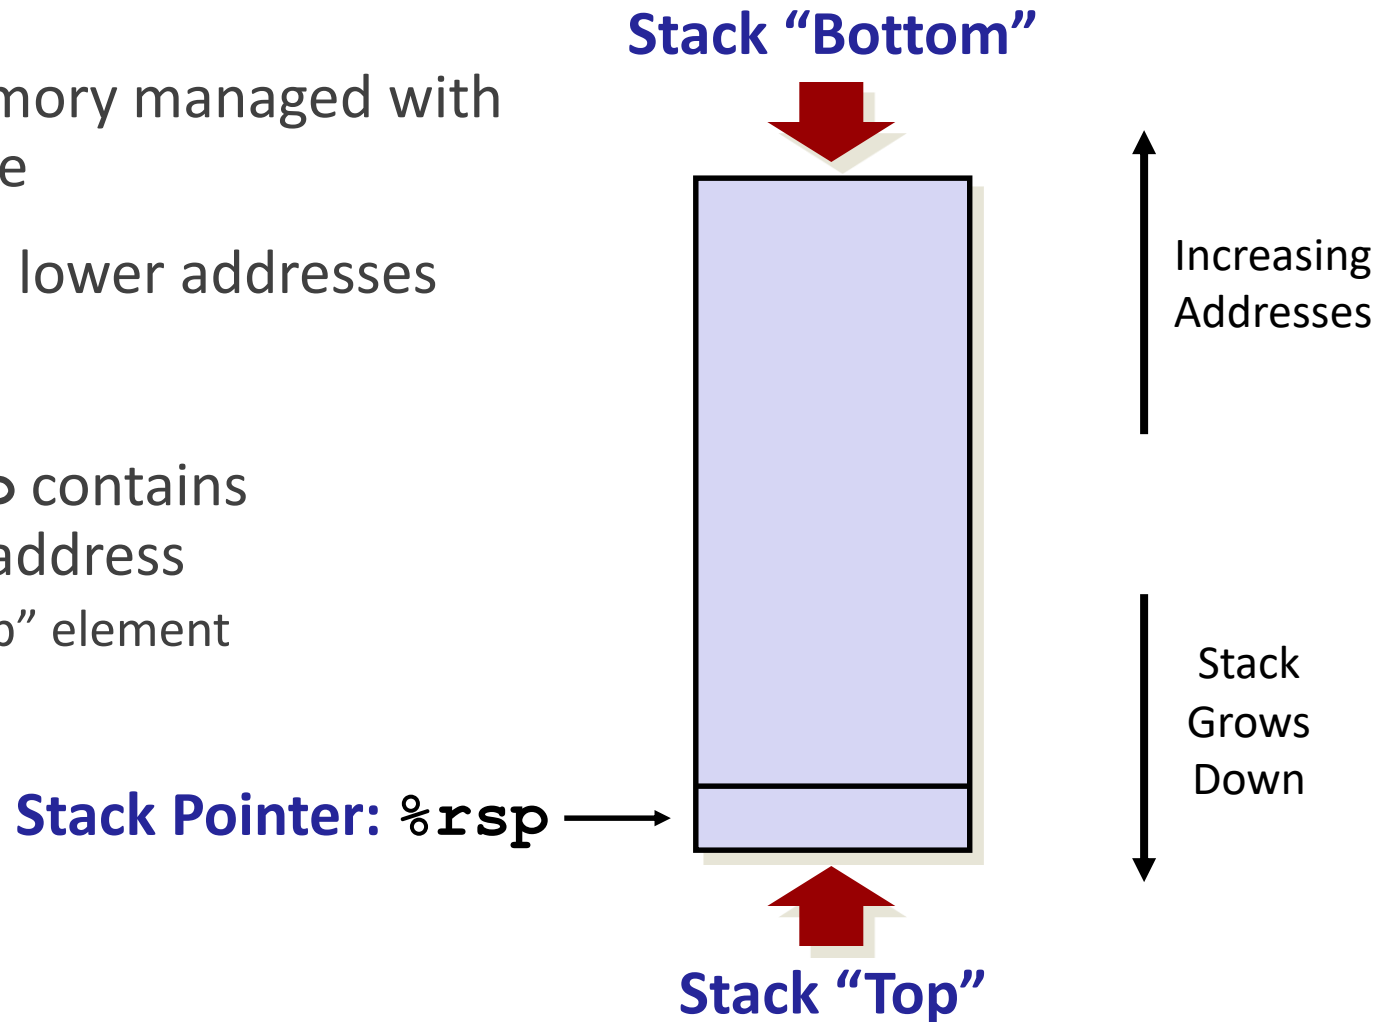
PORTLAND STATE UNIVERSITY

A solid green horizontal bar at the bottom of the slide.

Announcements

x86-64 Stack

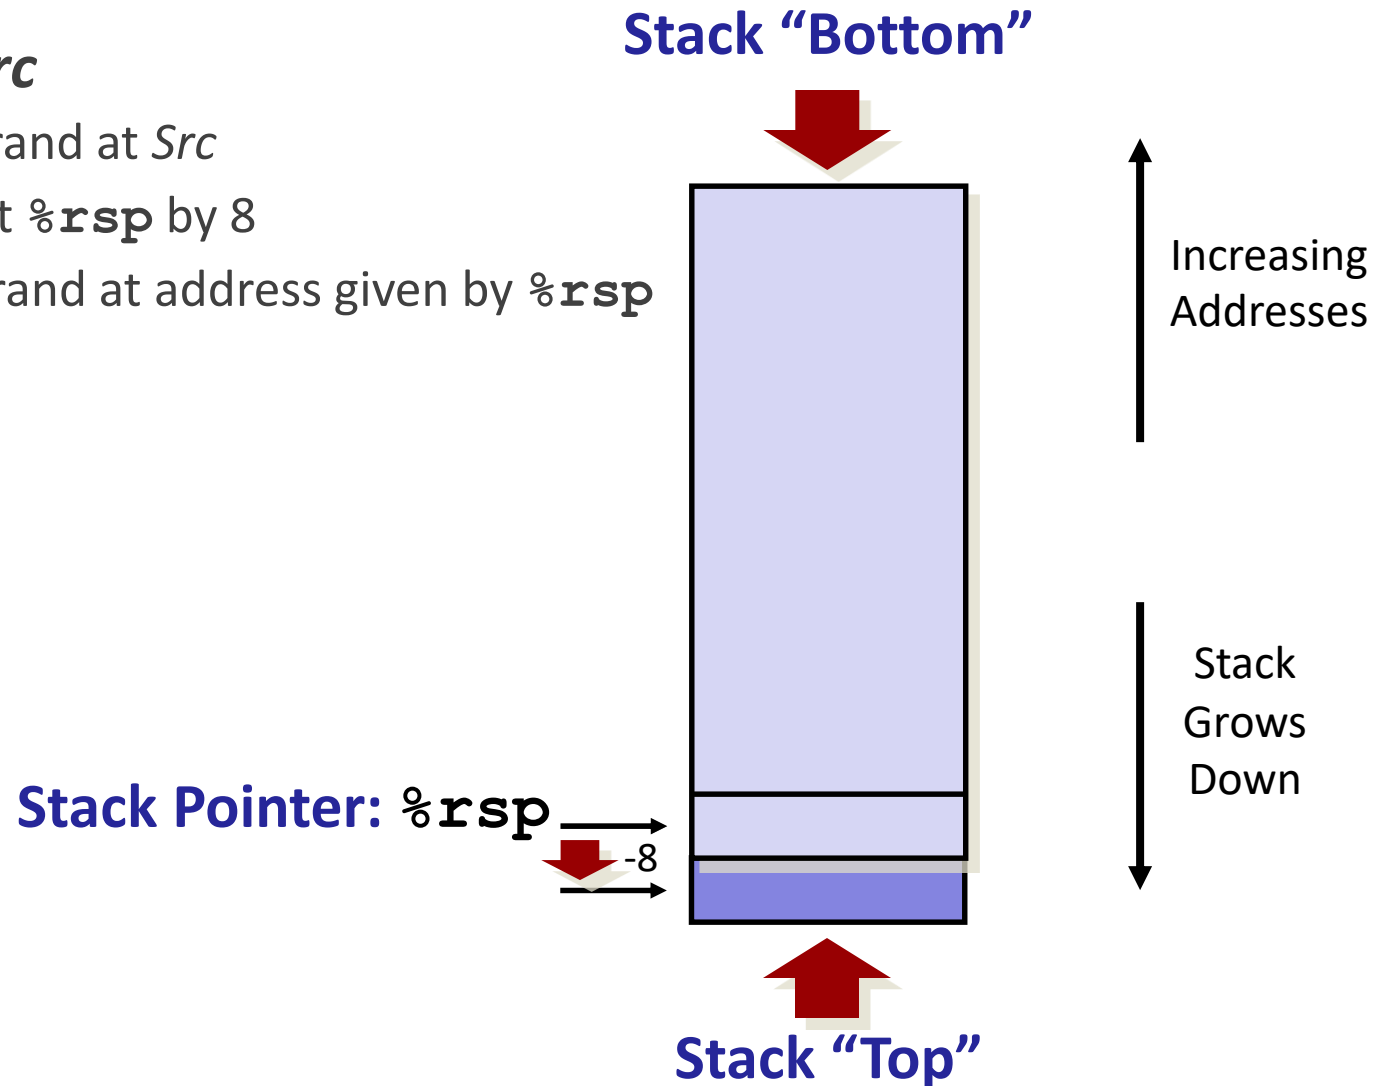
- Region of memory managed with stack discipline
- Grows toward lower addresses
- Register `%rsp` contains lowest stack address
 - address of “top” element



x86-64 Stack: Push

- **pushq *Src***

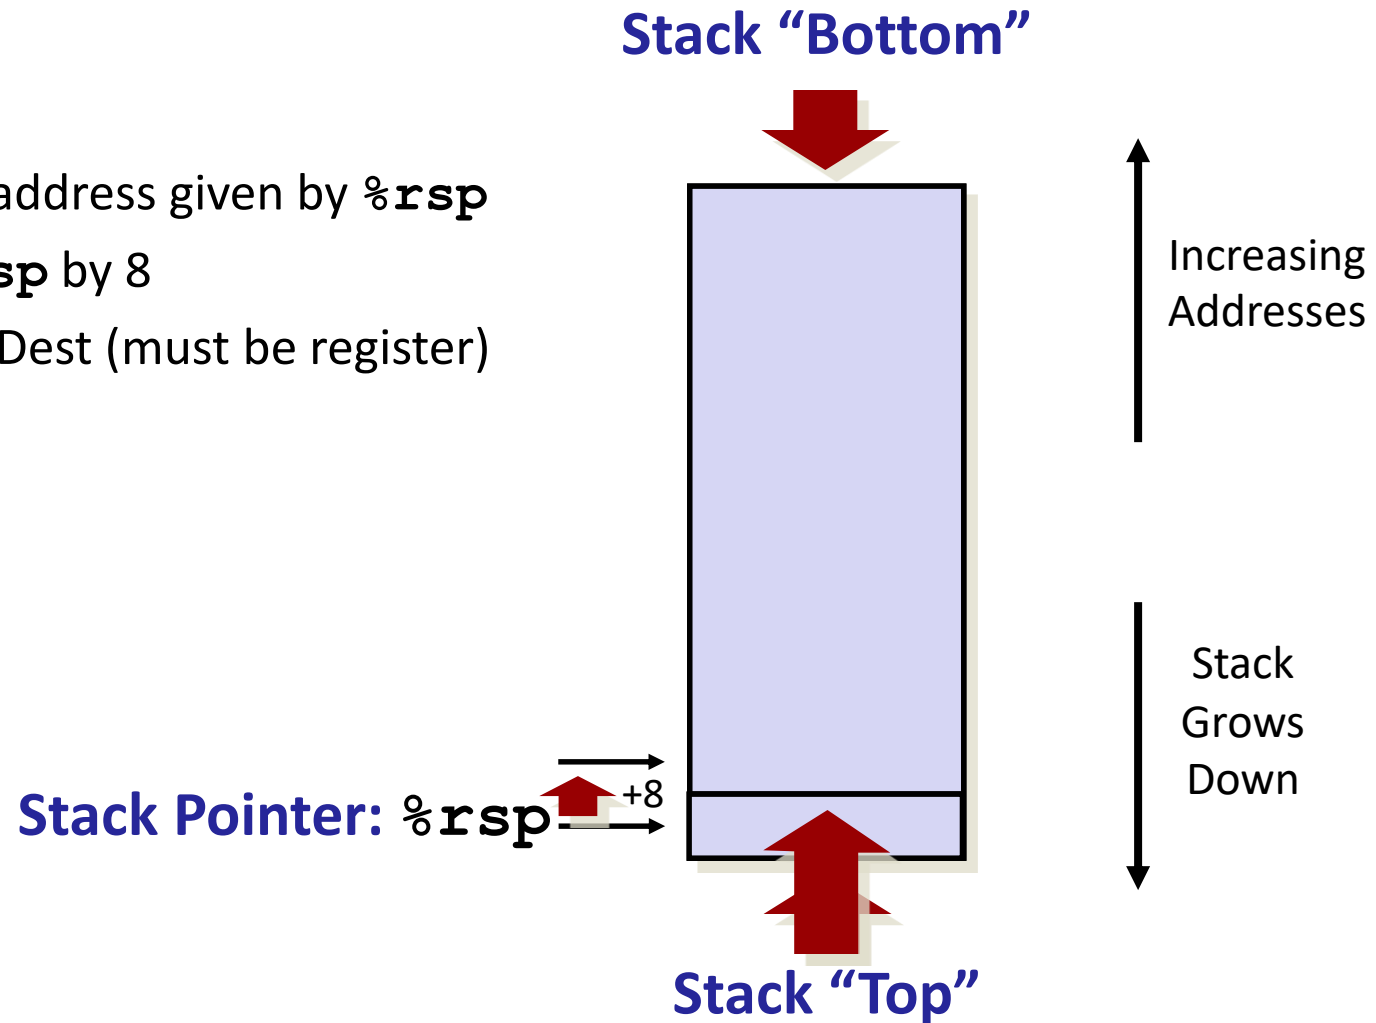
- Fetch operand at *Src*
- Decrement **%rsp** by 8
- Write operand at address given by **%rsp**



x86-64 Stack: Pop

■ `popq Dest`

- Read value at address given by `%rsp`
- Increment `%rsp` by 8
- Store value at `Dest` (must be register)



Procedure Control Flow

- Use stack to support procedure call and return
- Procedure call: `call label`
 - Push return address on stack
 - Jump to *label*
- Return address:
 - Address of the next instruction right after call
 - Example from disassembly
- Procedure return: `ret`
 - Pop address from stack
 - Jump to address

Procedure Control Flow Example

```
void multstore
(long x, long y, long *dest)
{
    long t = mult2(x, y);
    *dest = t;
}
```

```
00000000000400540 <multstore>:
400540: push    %rbx                # Save %rbx
400541: mov     %rdx,%rbx           # Save dest
400544: callq   400550 <mult2>      # mult2(x,y)
400549: mov     %rax, (%rbx)        # Save at dest
40054c: pop     %rbx                # Restore %rbx
40054d: retq                      # Return
```

```
long mult2
(long a, long b)
{
    long s = a * b;
    return s;
}
```

```
00000000000400550 <mult2>:
400550: mov     %rdi,%rax           # a
400553: imul    %rsi,%rax           # a * b
400557: retq                        # Return
```

Call Flow

```
00000000000400540 <multstore>:  
.  
.  
400544: callq 400550 <mult2>  
400549: mov  %rax, (%rbx)  
.  
.
```

```
00000000000400550 <mult2>:  
400550: mov  %rdi,%rax  
.  
.  
400557: retq
```

0x130

0x128

0x120

%rsp

0x120

%rip

0x400544

Call Flow

```
00000000000400540 <multstore>:  
.  
.  
400544: callq    400550 <mult2>  
400549: mov     %rax, (%rbx) ←  
.  
.
```

```
00000000000400550 <mult2>:  
400550: mov     %rdi, %rax ←  
.  
.  
400557: retq
```

0x130

0x128

0x120

0x118

%rsp

%rip

0x400549

0x118

0x400550

Call Flow

```
00000000000400540 <multstore>:  
.  
.  
400544: callq 400550 <mult2>  
400549: mov    %rax, (%rbx) ←  
.  
.
```

```
00000000000400550 <mult2>:  
400550: mov    %rdi,%rax  
.  
.  
400557: retq ←
```

0x130

0x128

0x120

0x118 0x400549

%rsp 0x118

%rip 0x400557

Call Flow

```
00000000000400540 <multstore>:  
.  
.  
400544: callq    400550 <mult2>  
400549: mov     %rax, (%rbx)  
.  
.
```

```
00000000000400550 <mult2>:  
400550: mov     %rdi, %rax  
.  
.  
400557: retq
```

0x130

0x128

0x120

%rsp

0x120

%rip

0x400549

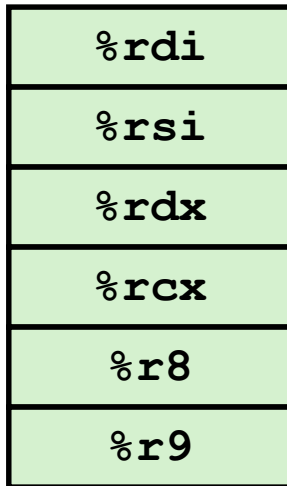
Calling Convention

- Specification detailing how a particular language and platform implement function calls
 - **Argument Passing**: How we pass arguments?
 - **Return Value**: How we return values?
 - **Register Saving Convention**: Which registers are preserved?
 - **Stack Frame Format and Management**: How stack is managed?
- Many Calling Conventions
 - CDECL – Used by C in x86 platforms
 - STDCALL – Windows API calls
 - **System V AMD64 – GCC/Linux in x64 platforms**

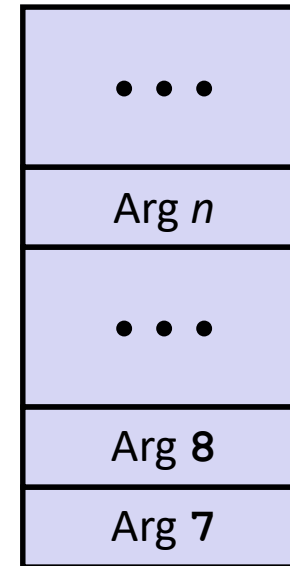
X64 Argument Passing

- Registers

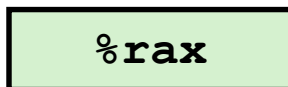
- First 6 arguments



- Stack



- Return value



- Only allocate stack space when needed

X64 Argument Passing

```
void multstore
(long x, long y, long *dest)
{
    long t = mult2(x, y);
    *dest = t;
}
```

```
0000000000400540 <multstore>:
    # x in %rdi, y in %rsi, dest in %rdx
    ...
400541: mov     %rdx,%rbx          # Save dest
400544: callq   400550 <mult2>     # mult2(x,y)
    # t in %rax
400549: mov     %rax, (%rbx)       # Save at dest
    ...
```

```
long mult2
(long a, long b)
{
    long s = a * b;
    return s;
}
```

```
0000000000400550 <mult2>:
    # a in %rdi, b in %rsi
400550: mov     %rdi,%rax          # a
400553: imul    %rsi,%rax          # a * b
    # s in %rax
400557: retq                                # Return
```

Register Saving Convention

- When procedure **yoo** calls **who**:
 - **yoo** is the *caller*
 - **who** is the *callee*
- Can register be used for temporary storage?

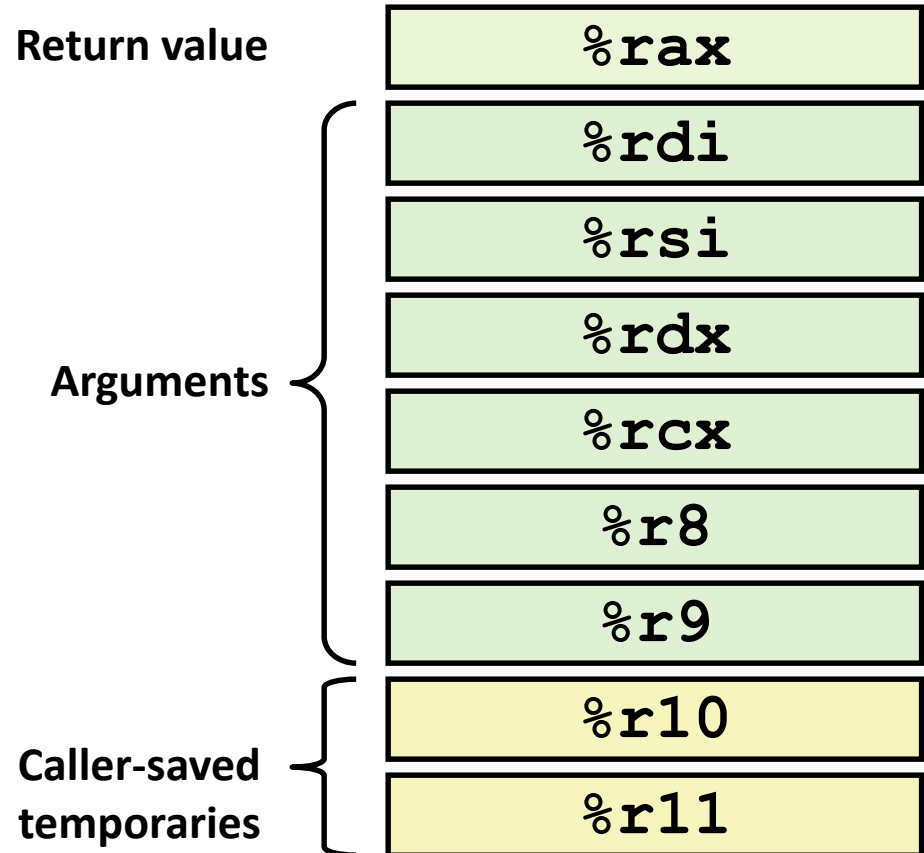
```
yoo:
    . . .
    movq $15213, %rdx
    call who
    addq %rdx, %rax
    . . .
    ret
```

```
who:
    . . .
    subq $18213, %rdx
    . . .
    ret
```

- Contents of register **%rdx** overwritten by **who**
- This could be trouble → something should be done!
 - Need some coordination

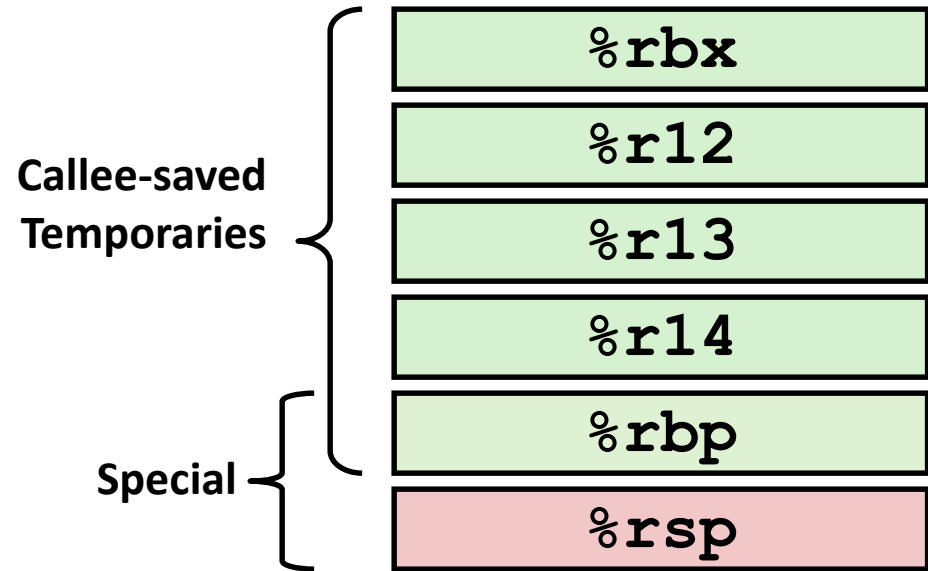
X64 Linux Register Saving

- **%rax**
 - Return value
 - Also caller-saved
 - Can be modified by procedure
- **%rdi, ..., %r9**
 - Arguments
 - Also caller-saved
 - Can be modified by procedure
- **%r10, %r11**
 - Caller-saved
 - Can be modified by procedure



X64 Linux Register Saving

- **%rbx, %r12, %r13, %r14**
 - Callee-saved
 - Callee must save & restore
- **%rbp**
 - Callee-saved
 - Callee must save & restore
 - May be used as frame pointer
 - Can mix & match
- **%rsp**
 - Special form of **callee save**
 - Restored to original value upon exit from procedure

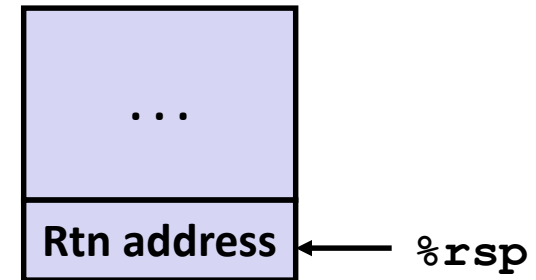


Register Saving Example

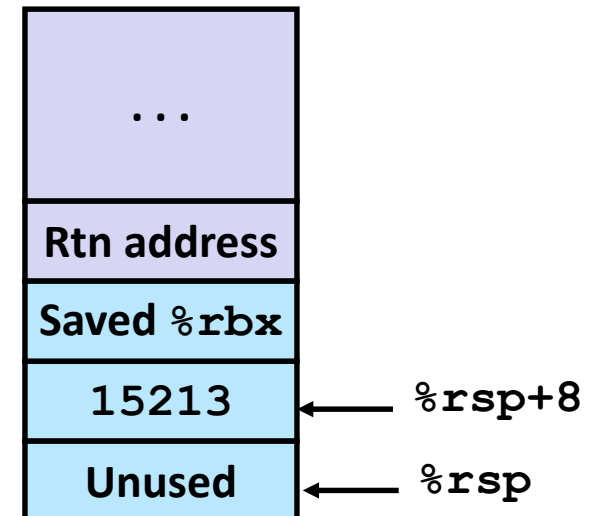
```
long call_incr2(long x) {  
    long v1 = 15213;  
    long v2 = incr(&v1, 3000);  
    return x+v2;  
}
```

```
call_incr2:  
    pushq    %rbx  
    subq     $16, %rsp  
    movq     %rdi, %rbx  
    movq     $15213, 8(%rsp)  
    movl     $3000, %esi  
    leaq     8(%rsp), %rdi  
    call     incr  
    addq     %rbx, %rax  
    addq     $16, %rsp  
    popq     %rbx  
    ret
```

Initial Stack Structure



Resulting Stack Structure

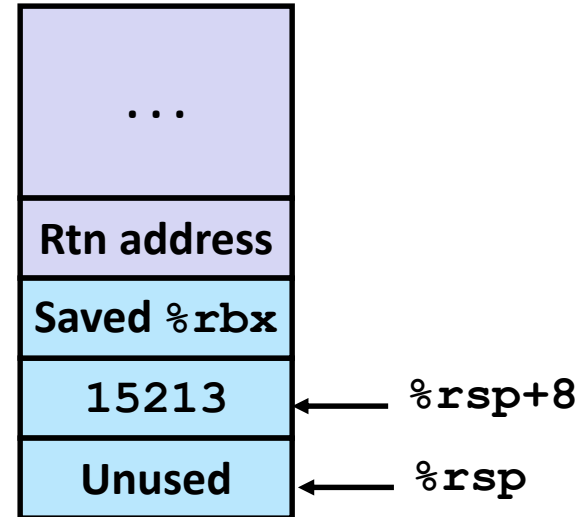


Register Saving Example

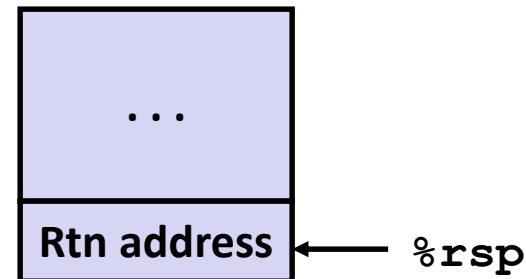
```
long call_incr2(long x) {  
    long v1 = 15213;  
    long v2 = incr(&v1, 3000);  
    return x+v2;  
}
```

```
call_incr2:  
    pushq    %rbx  
    subq     $16, %rsp  
    movq     %rdi, %rbx  
    movq     $15213, 8(%rsp)  
    movl     $3000, %esi  
    leaq     8(%rsp), %rdi  
    call     incr  
    addq     %rbx, %rax  
    addq     $16, %rsp  
    popq     %rbx  
    ret
```

Resulting Stack Structure



Pre-return Stack Structure



Why the Stack?

- Languages that support recursion
 - e.g., C, Pascal, Java
 - Code must be “*Reentrant*”
 - Multiple simultaneous instantiations of single procedure
 - Need some place to store state of each instantiation
 - Arguments
 - Local variables
 - Return pointer
- Stack discipline
 - State for given procedure needed for limited time
 - From when called to when return
 - Callee returns before caller does
- Stack allocated in ***Frames***
 - state for single procedure instantiation

Call Chain

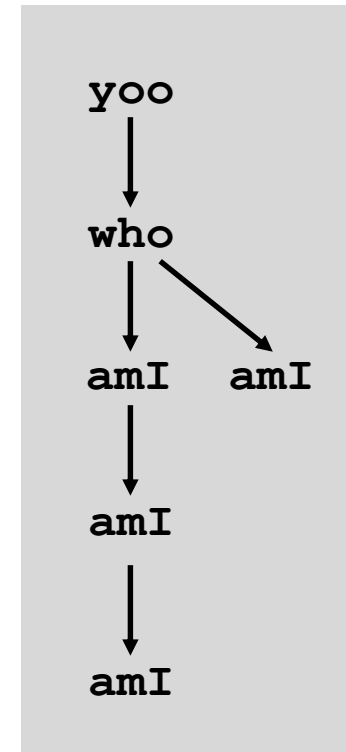
```
yoo (...)  
{  
  .  
  .  
  who ();  
  .  
  .  
}
```

```
who (...)  
{  
  . . .  
  amI ();  
  . . .  
  amI ();  
  . . .  
}
```

```
amI (...)  
{  
  .  
  .  
  amI ();  
  .  
  .  
}
```

Procedure `amI ()` is recursive

Example Call Chain



Stack Frame

- Contents

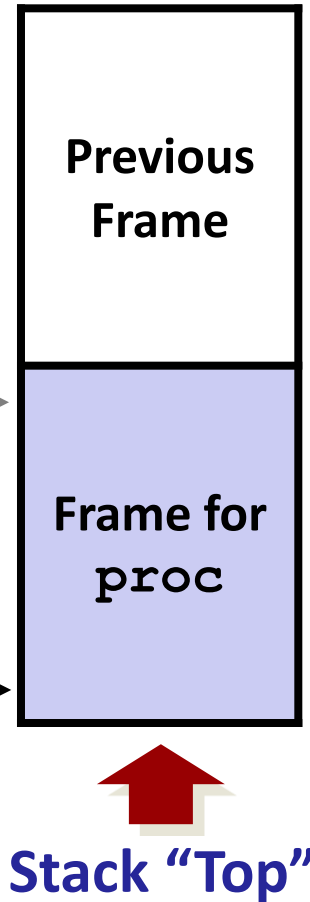
- Return information
- Local storage (if needed)
- Temporary space (if needed)

- Management

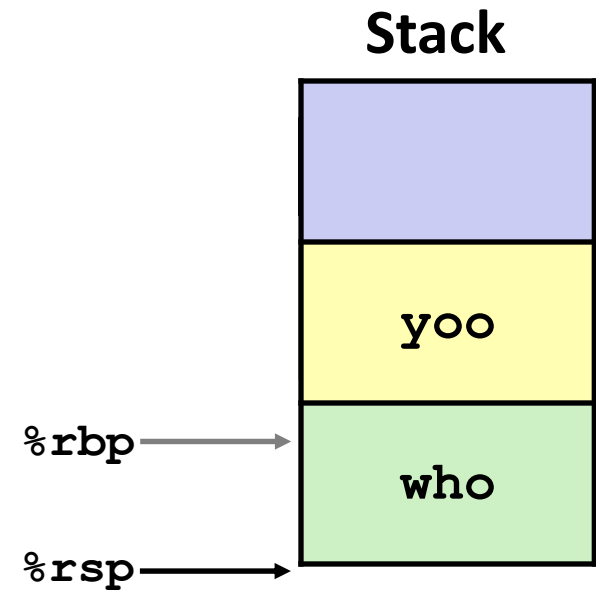
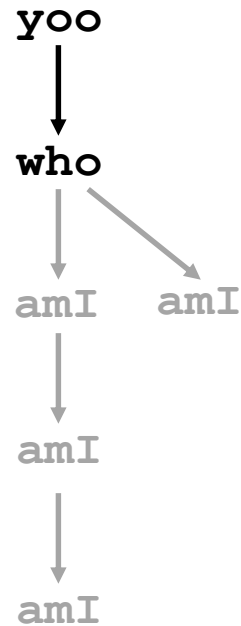
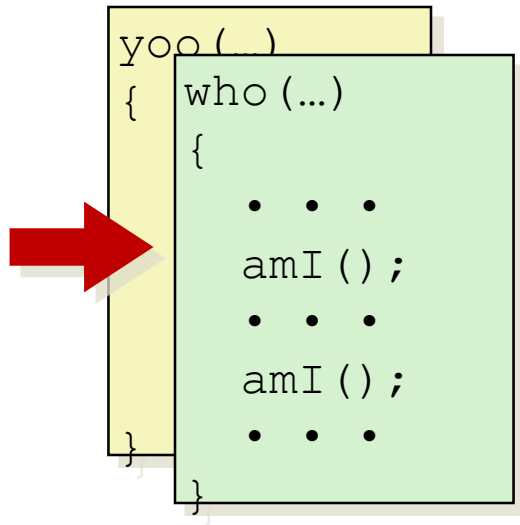
- Space allocated when enter procedure
 - “Set-up” code
 - Includes push by **call** instruction
- Deallocated when return
 - “Finish” code
 - Includes pop by **ret** instruction

Frame Pointer: **%rbp**
(Optional)

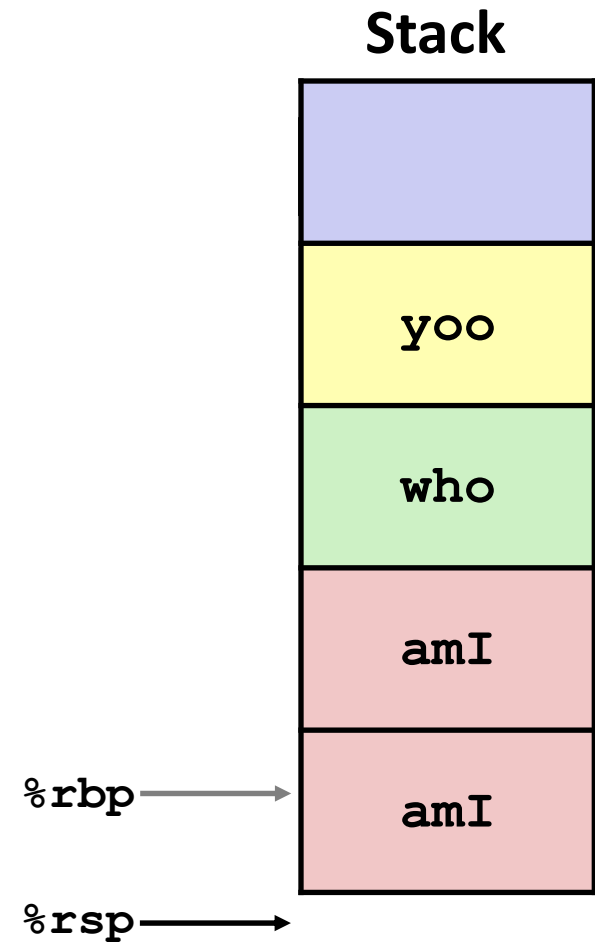
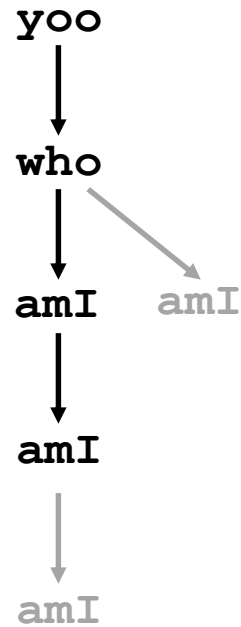
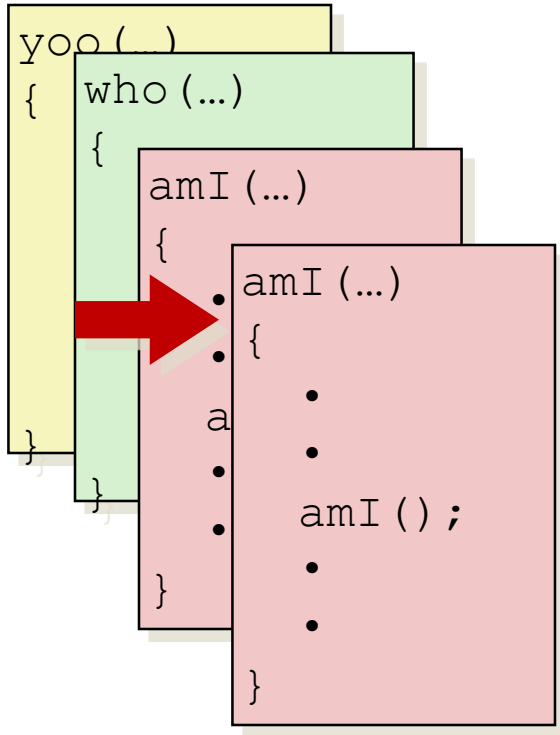
Stack Pointer: **%rsp**



Call Chain



Call Chain



Recursion

- Handled Without Special Consideration
 - Stack frames mean that each function call has private storage
 - Saved registers & local variables
 - Saved return pointer
 - Register saving conventions prevent one function call from corrupting another's data
 - Unless the C code explicitly does so (e.g., buffer overflow)
 - Stack discipline follows call / return pattern
 - If P calls Q, then Q returns before P
 - Last-In, First-Out
- Also works for mutual recursion
 - P calls Q; Q calls P

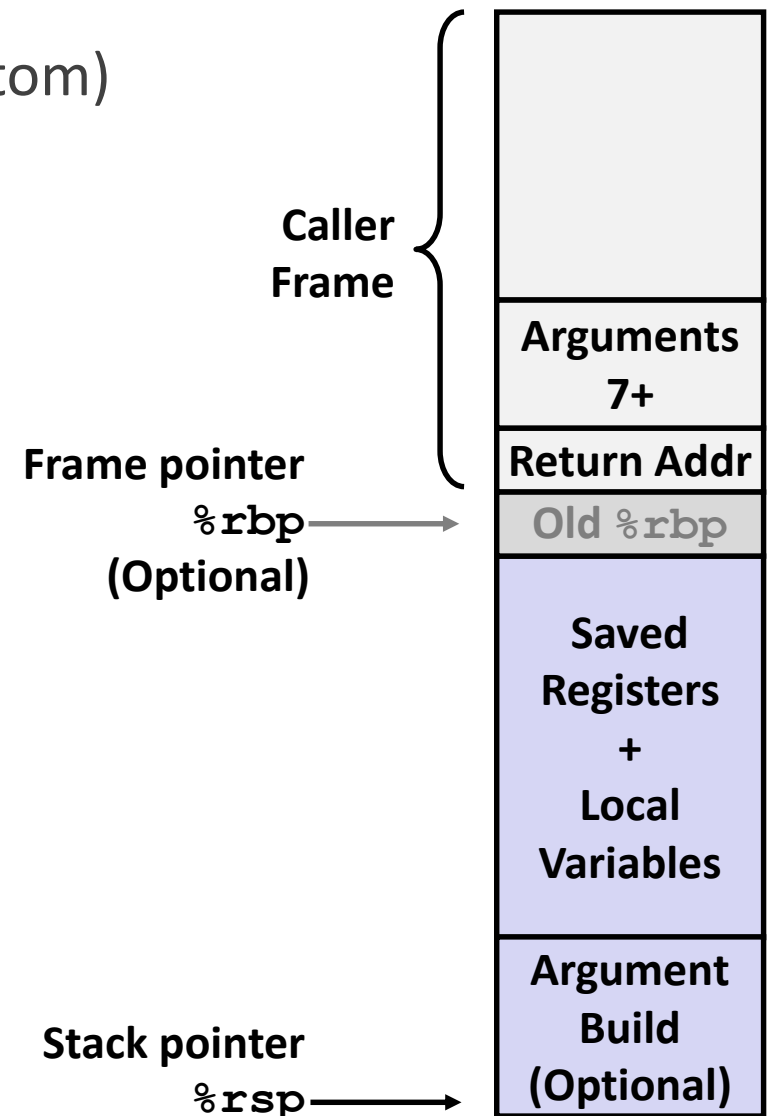
Linux Stack Frame

- Current Stack Frame (“Top” to Bottom)

- “Argument build:”
Parameters for function about to call
- Local variables
If can’t keep in registers
- Saved register context
- Old frame pointer (optional)

- Caller Stack Frame

- Return address
 - Pushed by `call` instruction
- Arguments for this call



GNU AS Assembler Directives

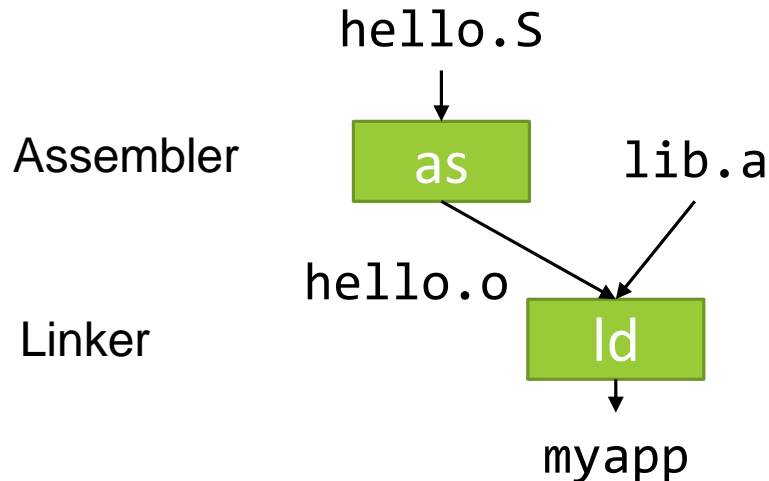
- Assembler directives are commands used as part of the assembler syntax.
- However, they are not part of the Intel x86-64 ISA.
 - Think of them as helper commands
- They start with the period character
- Some important Assembler directives are:
 - **.text**: Assemble the following statements onto the text section
 - **.data**: Assemble the following statements onto the data section
 - **.byte** *value*: Assemble a byte with the specified value
 - **.quad** *value*: Assemble a quadword with the specified value
 - **.asciz** *"string"*: Assemble a NULL terminated string
 - **.global** *symbol*: Makes the symbol visible to the linker (Global Symbol)
 - **.rept** *count*/**.endr**: Repeat the sequence of lines between the **.rept** directive and the next **.endr** directive count times
- More info: ftp://ftp.gnu.org/old-gnu/Manuals/gas/html_chapter/as_7.html

Assembly + libC “Hello World”

```
.global main                                     hello.S

.text
main:
    mov    $message, %rdi        # First parameter in %rdi
    call   puts                  # puts(message)
    ret                          # Return to C library code

message:
    .asciz "Hello, world"        # asciz puts a 0 byte at the end
```



```
as --64 hello.S -o hello.o
gcc -o myapp hello.o
```

Summary

- Stack is the right data structure for procedure call / return
 - If P calls Q, then Q returns before P: **Last-In First-Out (LIFO)**
- PUSH and POP instructions are used to control the stack
- CALL and RET are used to implement procedure calls
- Calling conventions are specifications about how a particular language and platform implement procedure calls
 - Argument Passing, Registers Saving Conventions, Stack Frame
- Recursion does not require any special handling