

CS201- Lecture 10

IA32 Data Access

RAOUL RIVAS

PORTLAND STATE UNIVERSITY

A solid green horizontal bar at the bottom of the slide.

Announcements


Array Allocation

- Basic Principle

T **A**[L];


- Array of data type T and length L
- Contiguously allocated region of $L * \text{sizeof}(T)$ bytes in memory

`char string[12];`




Remember x86 and x64
is byte addressable!


`int val[5];`




`double a[3];`



`char *p[3];`



Array Access Example



```
int get_digit
    (zip_dig z, int digit)
{
    return z[digit];
}
```

IA32

```
# %rdi = z
# %rsi = digit
movl (%rdi,%rsi,4), %eax # z[digit]
```

- Register `%rdi` contains starting address of array
- Register `%rsi` contains array index
- Desired digit at $\text{\%rdi} + 4 * \text{\%rsi}$
- Use memory reference $(\text{\%rdi}, \text{\%rsi}, 4)$

Array Loop Example

```
void zincr(zip_dig z) {  
    size_t i;  
    for (i = 0; i < ZLEN; i++)  
        z[i]++;  
}
```

```
# %rdi = z  
movl    $0, %eax           # i = 0  
jmp     .L3                # goto middle  
.L4:                        # loop:  
addl    $1, (%rdi,%rax,4)  # z[i]++  
addq    $1, %rax           # i++  
.L3:                        # middle  
cmpq    $4, %rax           # i:4  
jbe     .L4                # if <=, goto loop  
rep; ret
```

Static Multidimensional Arrays

- Declaration

$T \text{ } \mathbf{A}[R][C];$


- 2D array of data type T
- R rows, C columns
- Type T element requires K bytes

- Array Size

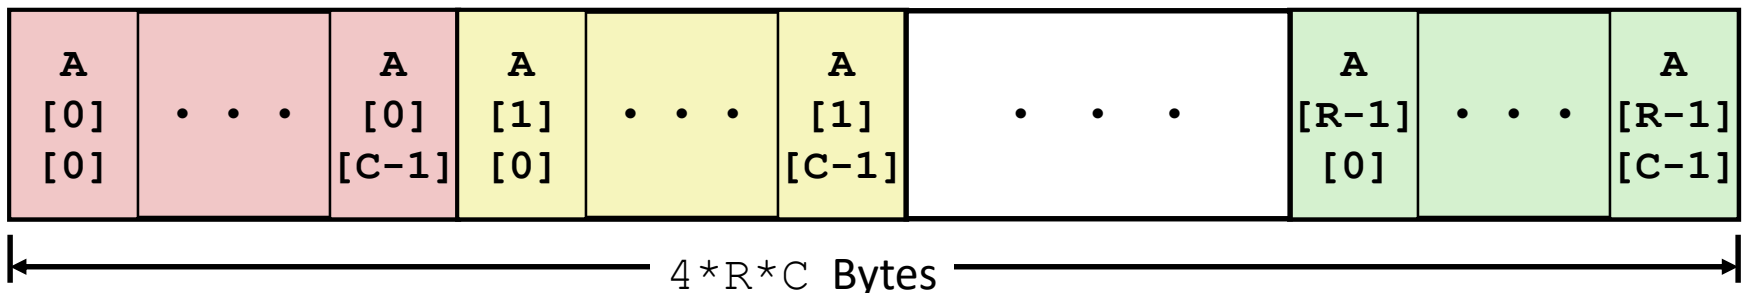
- $R * C * K$ bytes

- Arrangement

- Row-Major Ordering



`int A[R][C];`




Element Access (Static Array)


■ Array Elements

- $A[i][j]$ is element of type T , which requires K bytes
- Address $A + i * (C * K) + j * K = A + (i * C + j) * K$

```
int A[R][C];
```



Element Access (Static Array)



```
int get_pgh_digit
(int index, int dig)
{
    return pgh[index][dig];
}
```

```
leaq    (%rdi,%rdi,4), %rax    # 5*index
addl    %rax, %rsi             # 5*index+dig
movl    pgh(,%rsi,4), %eax     # M[pgh + 4*(5*index+dig)]
```

■ Array Elements


- `pgh[index][dig]` is `int`
- Address: `pgh + 20*index + 4*dig`
 - `= pgh + 4*(5*index + dig)`

Array of Pointer (Dynamic Array)

```
zip_dig cmu = { 1, 5, 2, 1, 3 };  
zip_dig mit = { 0, 2, 1, 3, 9 };  
zip_dig ucb = { 9, 4, 7, 2, 0 };
```


```
#define UCOUNT 3  
int *univ[UCOUNT] = {mit, cmu, ucb};
```

- Variable `univ` denotes array of 3 elements
- Each element is a pointer
 - 8 bytes
- Each pointer points to array of `int`'s



Element Access (Dynamic Array)

```
int get_univ_digit
(size_t index, size_t digit)
{
    return univ[index][digit];
}
```



```
salq    $2, %rsi          # 4*digit
addq     univ(,%rdi,8), %rsi # p = univ[index] + 4*digit
movl     (%rsi), %eax      # return *p
ret
```


■ Computation

- Element access **Mem[Mem[univ+8*index]+4*digit]**
- Must do two memory reads
 - First get pointer to row array
 - Then access element within array

Dynamic vs Static Array


Static array

```
int get_pgh_digit
(size_t index, size_t digit)
{
    return pgh[index][digit];
}
```



Dynamic array (Array of Pointers)

```
int get_univ_digit
(size_t index, size_t digit)
{
    return univ[index][digit];
}
```




Accesses looks similar in C, but address computations very different:

`Mem[pgh+20*index+4*digit]`

`Mem[Mem[univ+8*index]+4*digit]`

Structures


```
struct rec {  
    int a[4];  
    size_t i;  
    struct rec *next;  
};
```



- **Structure represented as block of memory**
 - Big enough to hold all of the fields
- **Fields ordered according to declaration**
 - Even if another ordering could yield a more compact representation
- **Compiler determines overall size + positions of fields**
 - Machine-level program has no understanding of the structures in the source code

Member Access

```
struct rec {  
    int a[4];  
    size_t i;  
    struct rec *next;  
};
```



■ Generating Pointer to Array Element


- Offset of each structure member determined at compile time
- Compute as $r + 4 \cdot idx$

```
int *get_ap  
(struct rec *r, size_t idx)  
{  
    return &r->a[idx];  
}
```

```
# r in %rdi, idx in %rsi  
leaq  (%rdi,%rsi,4), %rax  
ret
```

Data Alignment

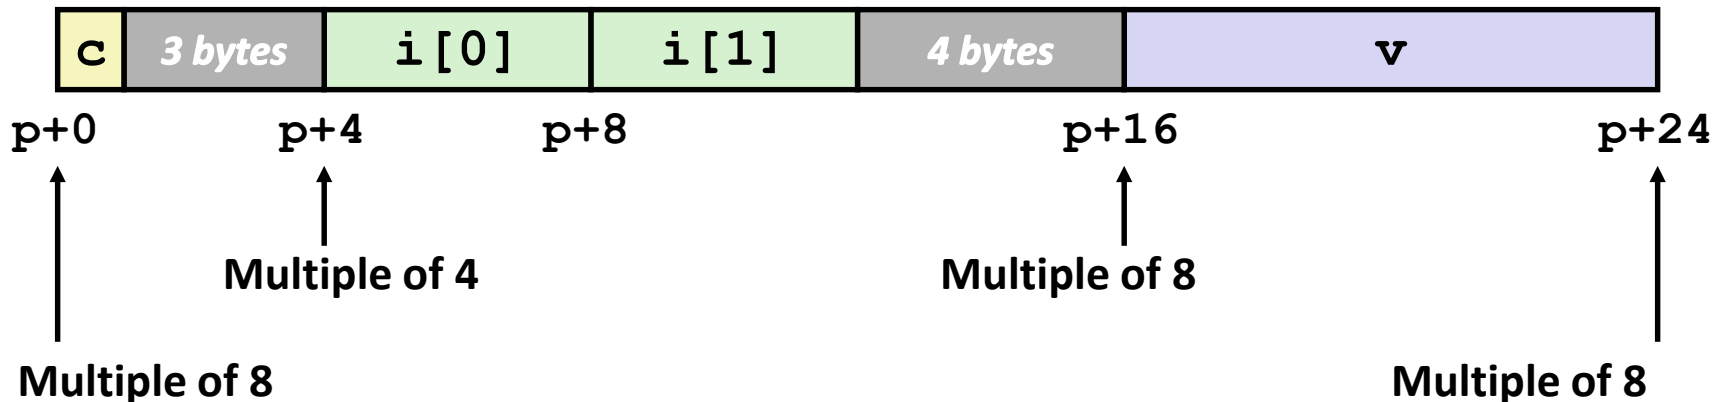
■ Unaligned Data



```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```

■ Aligned Data

- Primitive data type requires K bytes
- Address must be multiple of K



Alignment Principles

- Aligned Data
 - Primitive data type requires K bytes
 - Address must be multiple of K
 - Required on some machines; advised on x86-64
- Motivation for Aligning Data
 - Memory accessed by (aligned) chunks of 4 or 8 bytes (system dependent)
 - Inefficient to load or store datum that spans quad word boundaries
 - Virtual memory trickier when datum spans 2 pages
- Compiler
 - Inserts gaps in structure to ensure correct alignment of fields

X64 Alignment

- 1 byte: **char**
 - no restrictions on address
- 2 bytes: **short**
 - lowest 1 bit of address must be 0_2
- 4 bytes: **int, float**
 - lowest 2 bits of address must be 00_2
- 8 bytes: **double, long, char ***
 - lowest 3 bits of address must be 000_2
- 16 bytes: **long double** (GCC on Linux)
 - lowest 4 bits of address must be 0000_2

Structure Alignment

■ Within structure:

- Must satisfy each element's alignment requirement

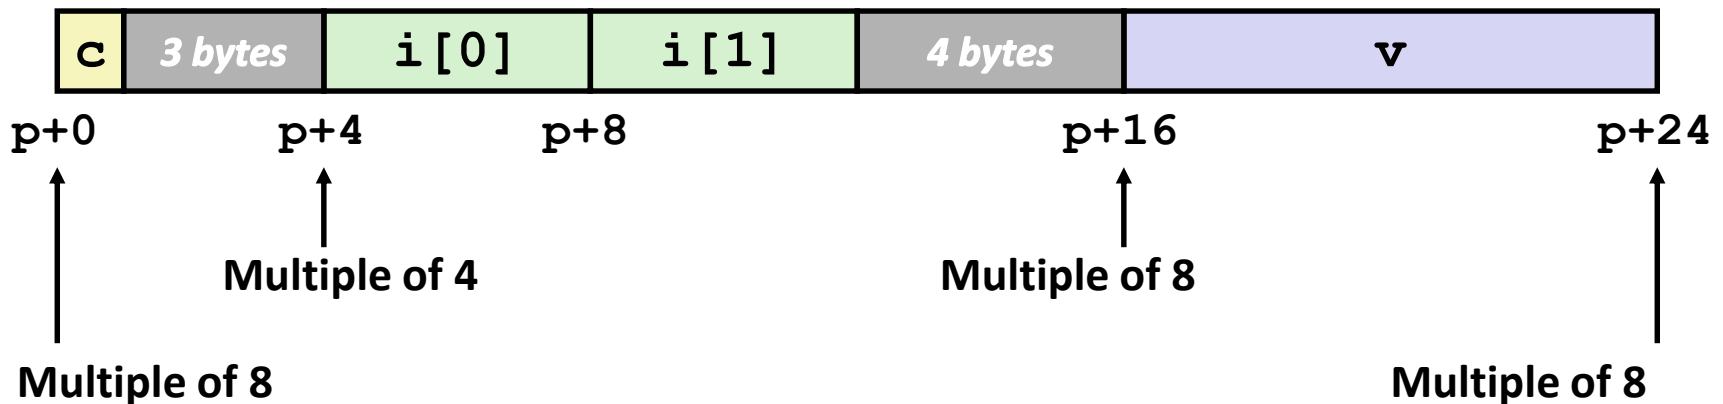
■ Overall structure placement

- Each structure has alignment requirement **K**
 - **K** = Largest alignment of any element
- Initial address & structure length must be multiples of **K**

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```

■ Example:

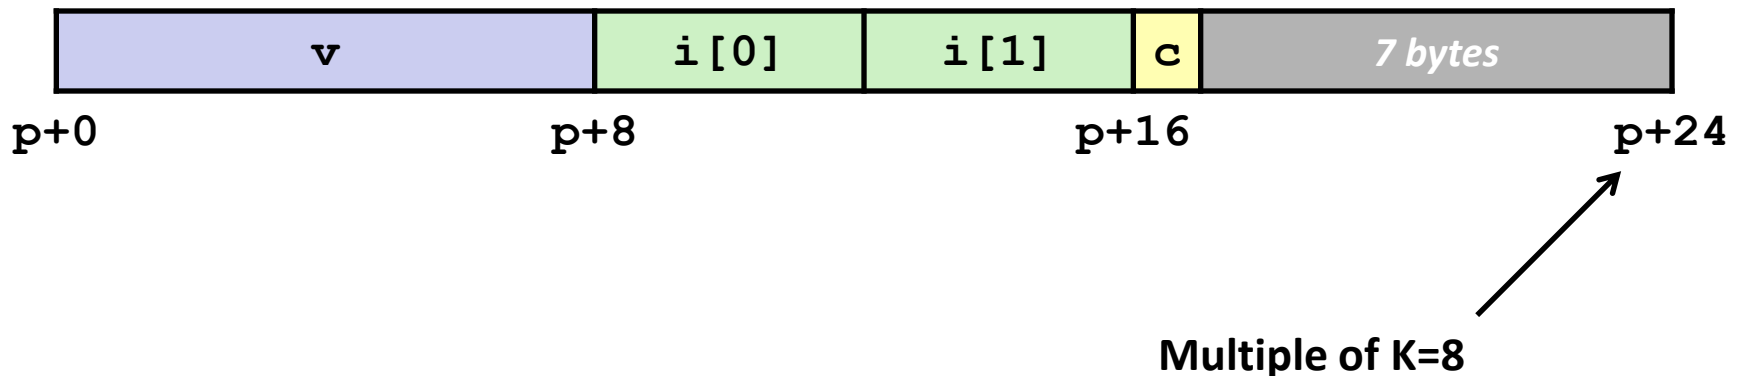
- **K** = 8, due to **double** element



Structure Alignment

- For largest alignment requirement K
- Overall structure must be multiple of K


```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} *p;
```



Arrays of Structures

- Overall structure length multiple of K
- Satisfy alignment requirement for every element

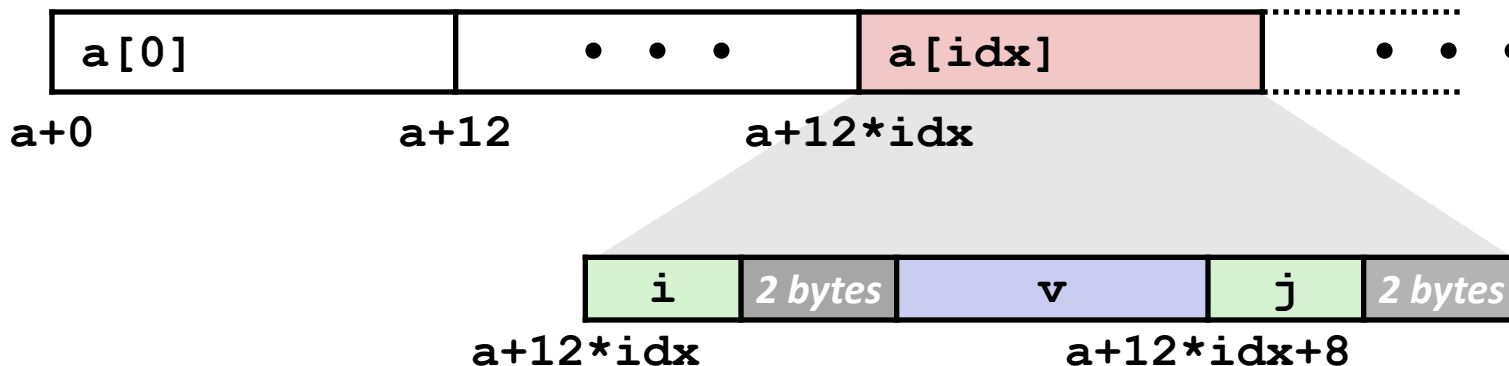
```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} a[10];
```



Accessing Arrays of Structures

- **Compute array offset $12 \cdot \text{idx}$**
 - `sizeof(S3)`, including alignment spacers
- **Element j is at offset 8 within structure**
- **Assembler gives offset $a+8$**
 - Resolved during linking

```
struct S3 {  
    short i;  
    float v;  
    short j;  
} a[10];
```




```
short get_j(int idx)  
{  
    return a[idx].j;  
}
```

```
# %rdi = idx  
leaq (%rdi,%rdi,2),%rax # 3*idx  
movzwl a+8(,%rax,4),%eax
```

Structures Alignment Optimization

■ Put large data types first

```
struct S4 {  
    char c;  
    int i;  
    char d;  
} *p;
```



```
struct S5 {  
    int i;  
    char c;  
    char d;  
} *p;
```

■ Result




Data Serialization

```
struct S4 {  
    char c;  
    int i;  
    char d;  
} *p;
```



Data Exchange





Internet
Flash Drives

```
struct S4 {  
    char c;  
    int i;  
    char d;  
} *p;
```



Data Corruption




Never send a raw structure across the Network

Data Serialization

- If the language includes serialization routines use them
 - Java and C#
- If not copy structure member by member to a contiguous buffer before send
- Endianness – Use common agreement of formats
 - hton macros in C
- Data sizes – Use independent data types
 - StdInt.h – uint16_t, uint32_t, etc
- Some libraries allow for data serialization trivially
 - Google Protocol Buffers Library

Data Serialization



```
char b[sizeof(char)*2+sizeof(int)];   Serialize.c
int off=0;
memcpy(b, &p->c, sizeof(char));
off = sizeof(char);
memcpy(b + off, &p->i, sizeof(int));
off += sizeof(int);
memcpy(b + off, &p->d, sizeof(char));
send(b, sizeof(b));
```


```
struct S4 {
    char c;
    int i;
    char d;
} *p;
```


Linked Lists

■ C Code

```
void set_val
(struct rec *r, int val)
{
    while (r) {
        int i = r->i;
        r->a[i] = val;
        r = r->next;
    }
}
```

```
struct rec {
    int a[4];
    int i;
    struct rec *next;
};
```



Element i

Register	Value
%rdi	r
%rsi	val

```
.L11:                                # loop:
    movslq    16(%rdi), %rax          # i = M[r+16]
    movl      %esi, (%rdi,%rax,4)     # M[r+4*i] = val
    movq      24(%rdi), %rdi          # r = M[r+24]
    testq     %rdi, %rdi              # Test r
    jne       .L11                   # if !=0 goto loop
```

Unions

- Allocate according to largest element
- Can only use one field at a time


```
union U1 {  
    char c;  
    int i[2];  
    double v;  
} *up;
```

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *sp;
```


Any real life example?

Java Number class implementation

Union



Structure



Summary

- Static and Dynamic array access in C use the same semantics (operator []) but address computation is different
- Members of Structures must be properly aligned to avoid performance penalty
- Structures must be serialized before saving them to files or before sending them across the network
- Unions allow a single variable to represent multiple types